

# Verifying concurrent, crash-safe systems with **Perennial**

**Tej Chajed**, Joseph Tassarotti\*, Frans Kaashoek, Nickolai Zeldovich

MIT and \*Boston College


# Many systems need concurrency and crash safety

Examples: file systems, databases, and key-value stores


Make strong guarantees about keeping your data safe

Achieve high performance with concurrency


# Simple example: replicated disk



# Simple example: replicated disk



# Simple example: replicated disk




# Replicated disk is subtle

```
func write(a: addr, v: block) {  
    lock_address(a)  
    d1.write(a, v)  
    d2.write(a, v)  
    unlock_address(a)  
}
```

# Replicated disk is subtle


```
func write(a: addr, v: block) {  
    lock_address(a)  
    d1.write(a, v)  
    d2.write(a, v)  
    unlock_address(a)  
}
```



**what if system crashes here?**  
**what if disk 1 fails?**

# Replicated disk is subtle

```
func write(a: addr, v: block) {  
    lock_address(a)  
    d1.write(a, v)  
    d2.write(a, v)  
    unlock_address(a)  
}
```




what if system crashes here?  
what if disk 1 fails?

```
// runs on reboot  
func recover() {  
    for a in ... {  
        // copy from d1 to d2  
    }  
}
```



# Replicated disk is subtle

```
func write(a: addr, v: block) {  
    lock_address(a)  
    d1.write(a, v)  
    d2.write(a, v)  
    unlock_address(a)  
}
```

 **what if system crashes here?**  
**what if disk 1 fails?**

```
// runs on reboot  
func recover() {  
    for a in ... {  
        // copy from d1 to d2  
    }  
}
```

```
func read(a: addr): block {  
    lock_address(a)  
    v, ok := d1.read(a)  
    if !ok {  
        v, _ = d2.read(a)  
    }  
    unlock_address(a)  
    return v  
}
```

**Goal: systematically reason about all executions  
with formal verification**

# Existing verification frameworks do not support concurrency and crash safety

verified crash safety

FSCQ [SOSP '15]

Yggdrasil [OSDI '16]

DFSCQ [SOSP '17]

...

verified concurrency

CertiKOS [OSDI '16]

CSPEC [OSDI '18]

AtomFS [SOSP '19]

...

**no system can do both**

# Combining verified crash safety and concurrency is challenging

Crash and recovery can interrupt a critical section

Crash wipes in-memory state

Recovery logically completes crashed threads' operations

# Perennial's techniques address challenges integrating crash safety into concurrency reasoning

Crash and recovery can interrupt a critical section

➡ **leases**

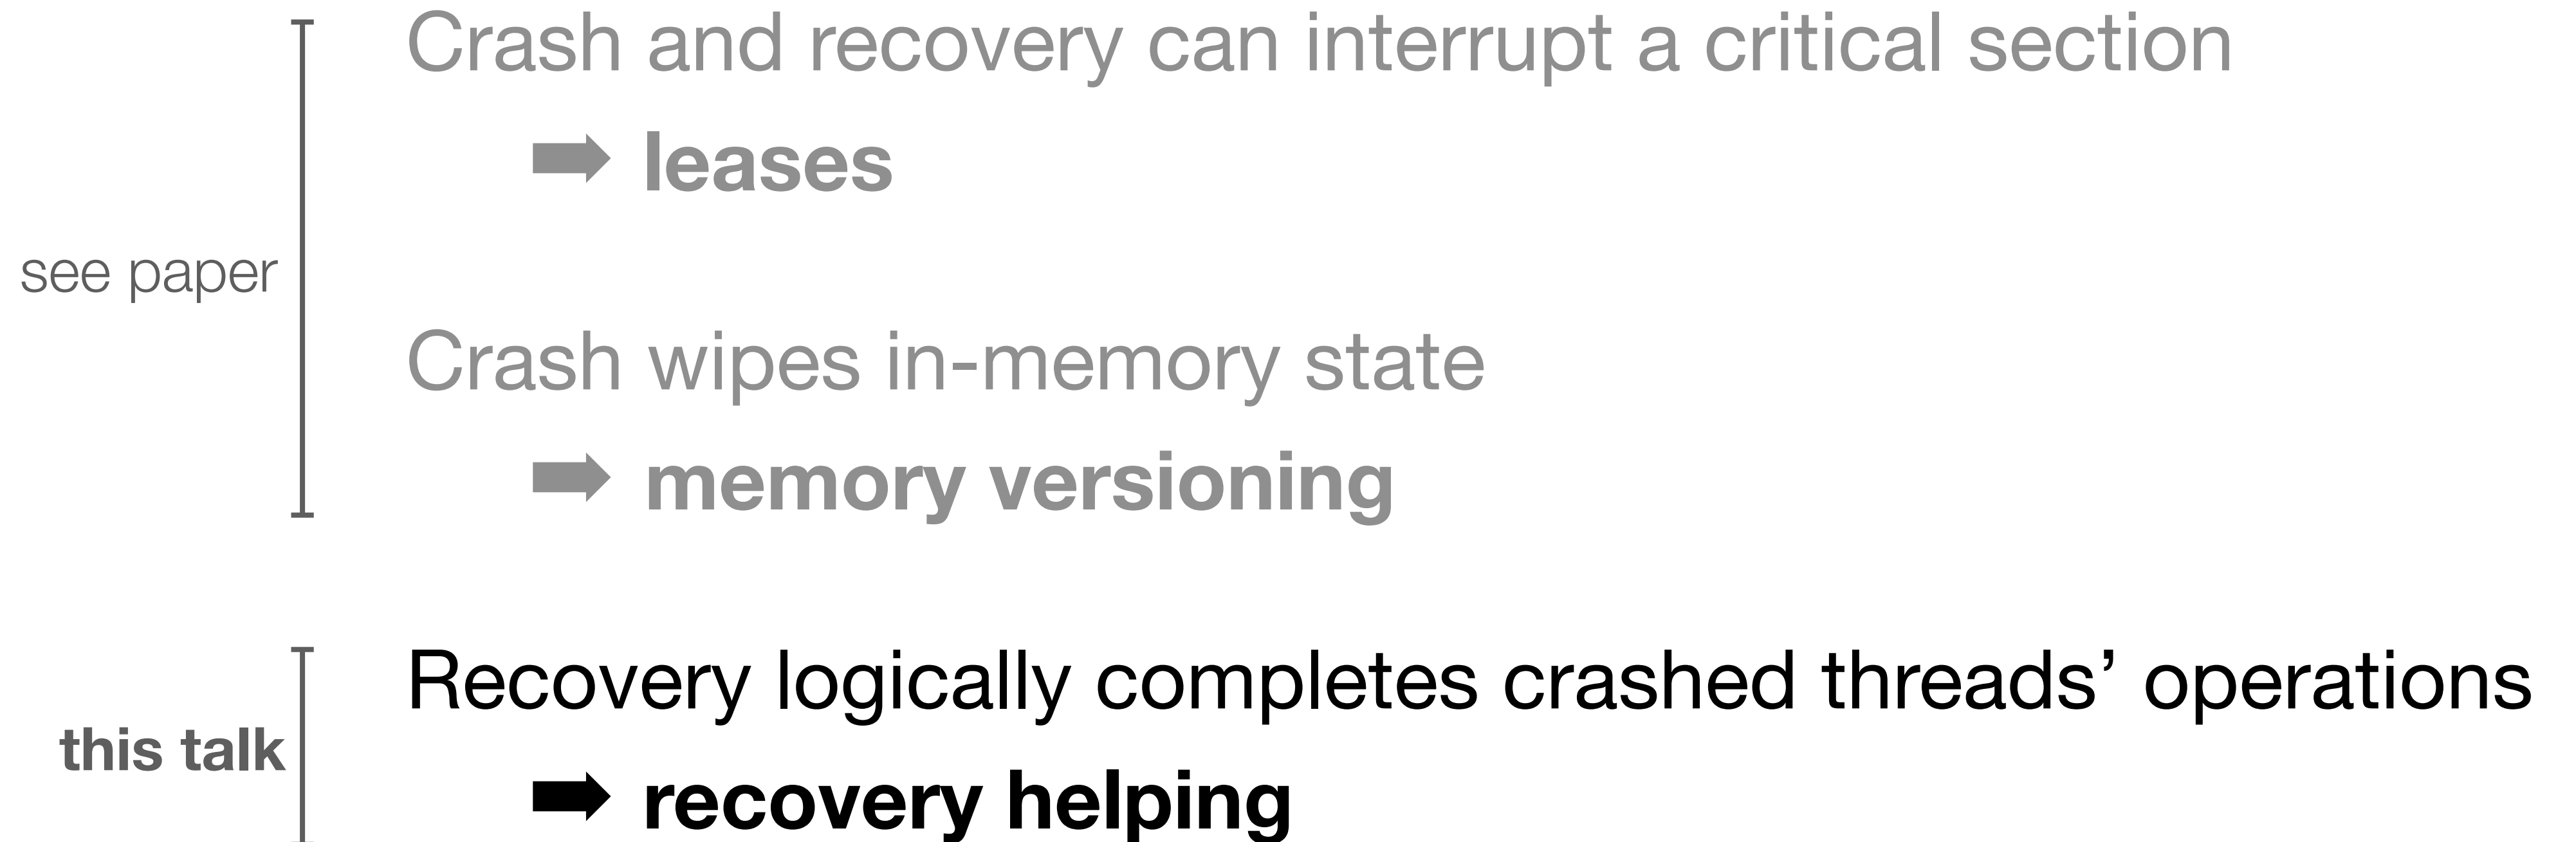
Crash wipes in-memory state

➡ **memory versioning**

Recovery logically completes crashed threads' operations

➡ **recovery helping**

# Perennial's techniques address challenges integrating crash safety into concurrency reasoning



# Contributions

Perennial: framework for reasoning about crashes and concurrency

see paper **Goose: reasoning about Go implementations**

Evaluation: verified mail server written in Go with Perennial

# Specifying correctness: concurrent recovery refinement

All operations are **correct** and **atomic wrt  
concurrency and crashes**


Recovery repairs system after reboot




# Proving the replicated disk correct

# Proving refinement with forward simulation: relate code and spec states

Background




# Proving refinement with forward simulation: prove every operation has a commit point



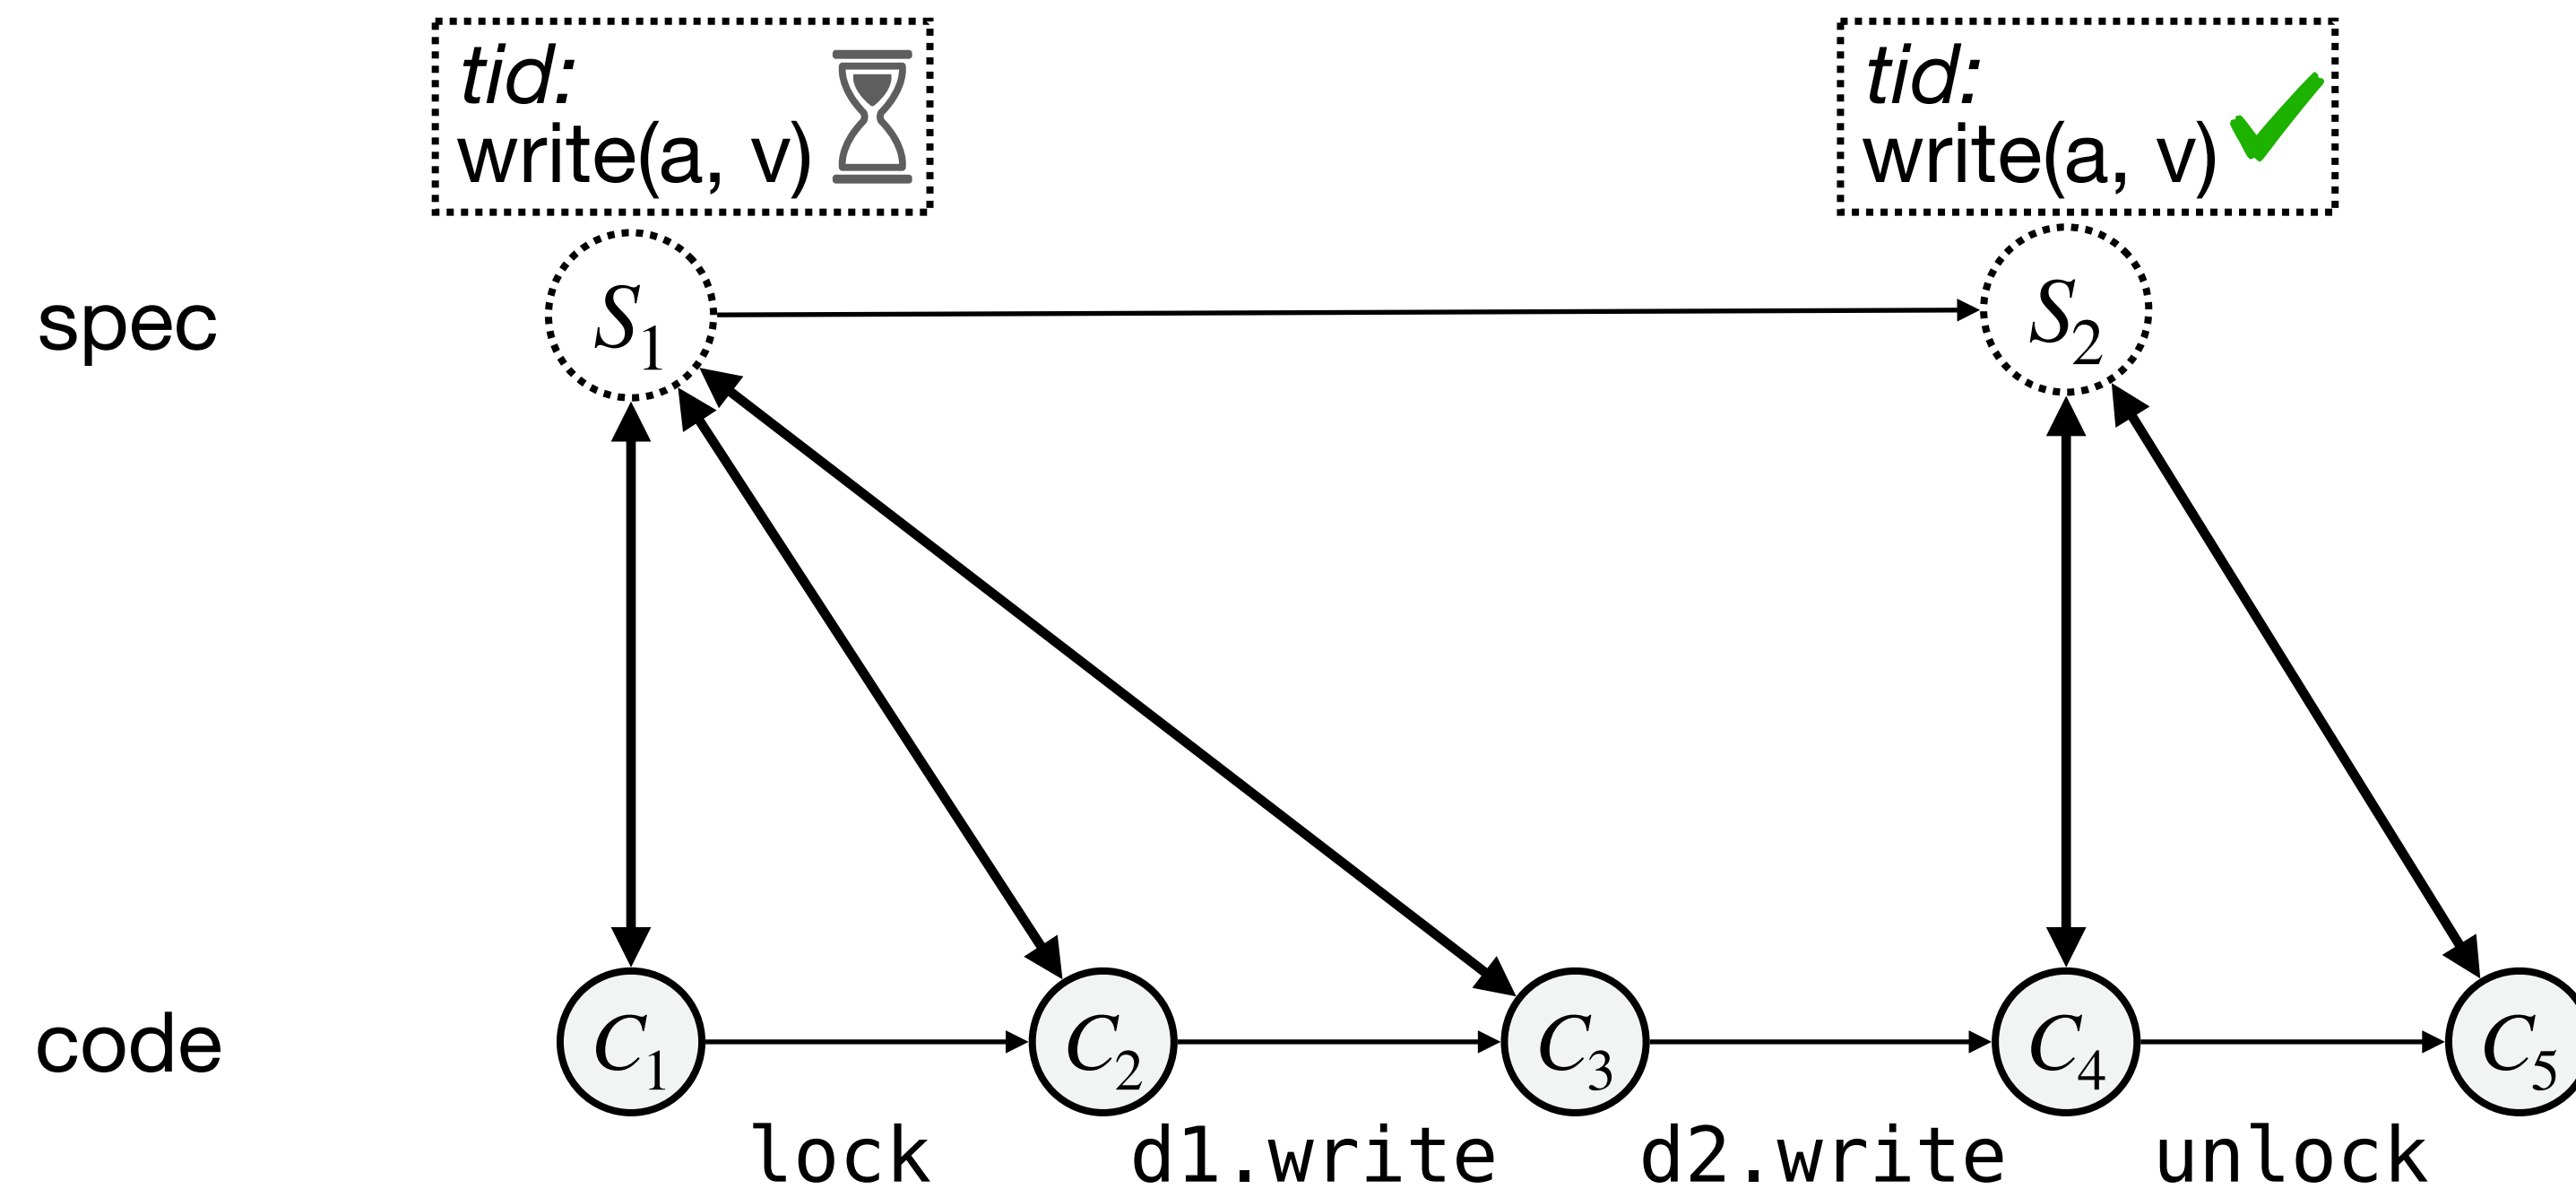
1. Write down abstraction relation between code and spec states

# Proving refinement with forward simulation: prove every operation has a commit point




1. Write down abstraction relation between code and spec states
2. Prove every operation commits

# Proving refinement with forward simulation: prove every operation has a commit point



1. Write down abstraction relation between code and spec states
2. Prove every operation commits
3. Prove abstraction relation is preserved

# Abstraction relation for the replicated disk



abstraction relation:

$$\text{!locked}(a) \implies \begin{aligned} &\sigma[a] = d_1[a] \\ &\wedge \sigma[a] = d_2[a] \end{aligned}$$

(if the disk has not failed)

# Crashing breaks the abstraction relation

```
func write(a: addr,  
          v: block) {
```


```
    lock_address(a)  
    d1.write(a, v)
```

lock reverts to being free,  
but disks are not in-sync

abstraction relation:

$$\text{!locked}(a) \implies \begin{array}{l} \sigma[a] = d_1[a] \\ \wedge \sigma[a] = d_2[a] \end{array}$$


# So far: abstraction relation always holds




$R$  abstraction relation



# Separate a **crash invariant** from the abstraction relation



# Recovery proof uses the crash invariant to restore the abstraction relation



# Proving recovery correct: makes writes atomic

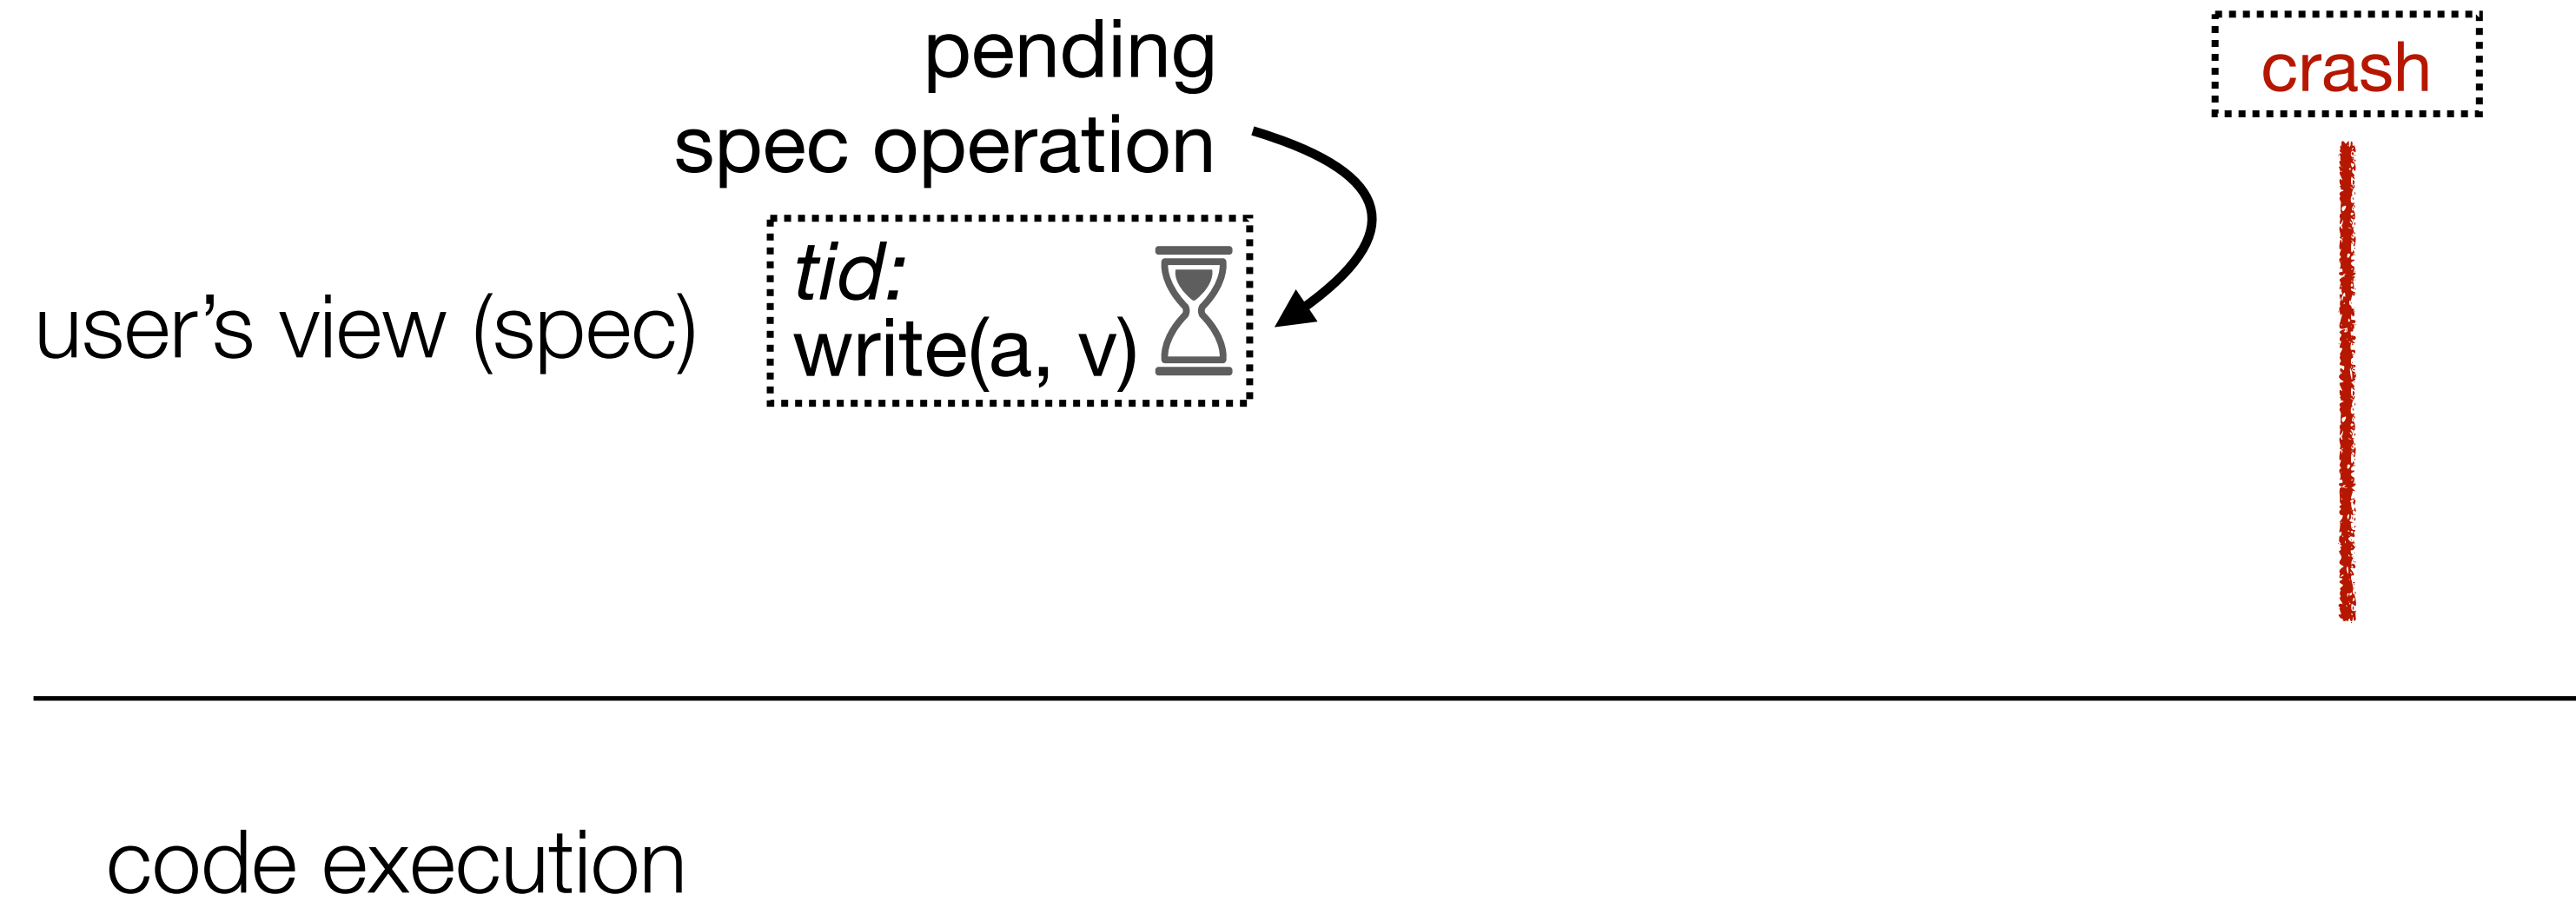
```
func write(a: addr,  
          v: block) {
```

```
    lock_address(a)  
    d1.write(a, v)
```

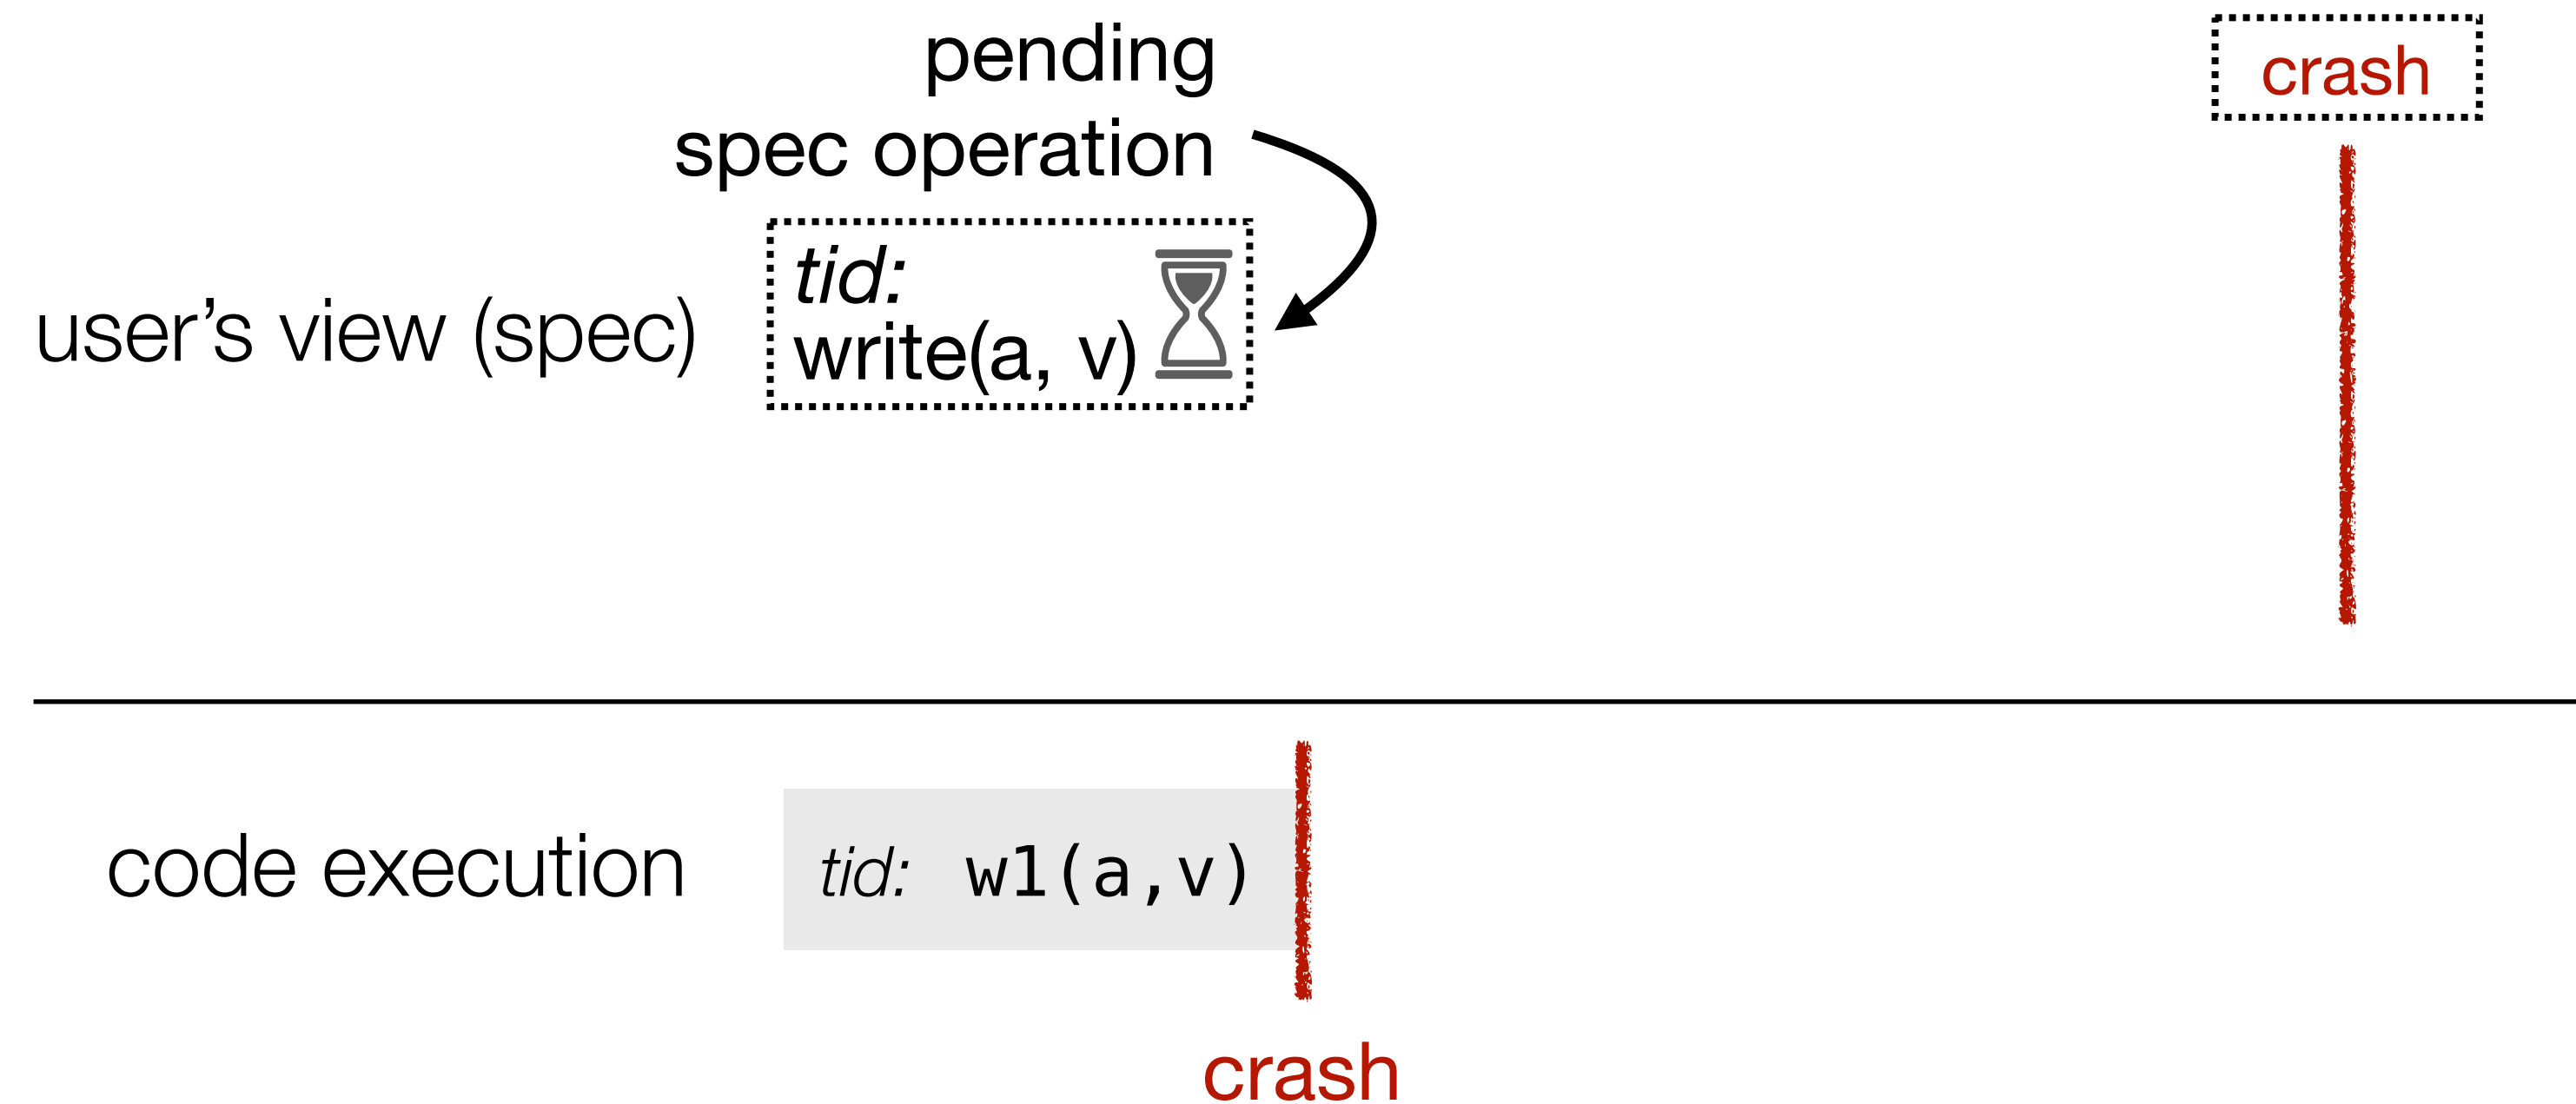
---

```
func recover() {  
    for a in ... {  
        v, ok := d1.read(a)  
        if !ok { ... }  
        d2.write(a, v)  
    }  
}
```

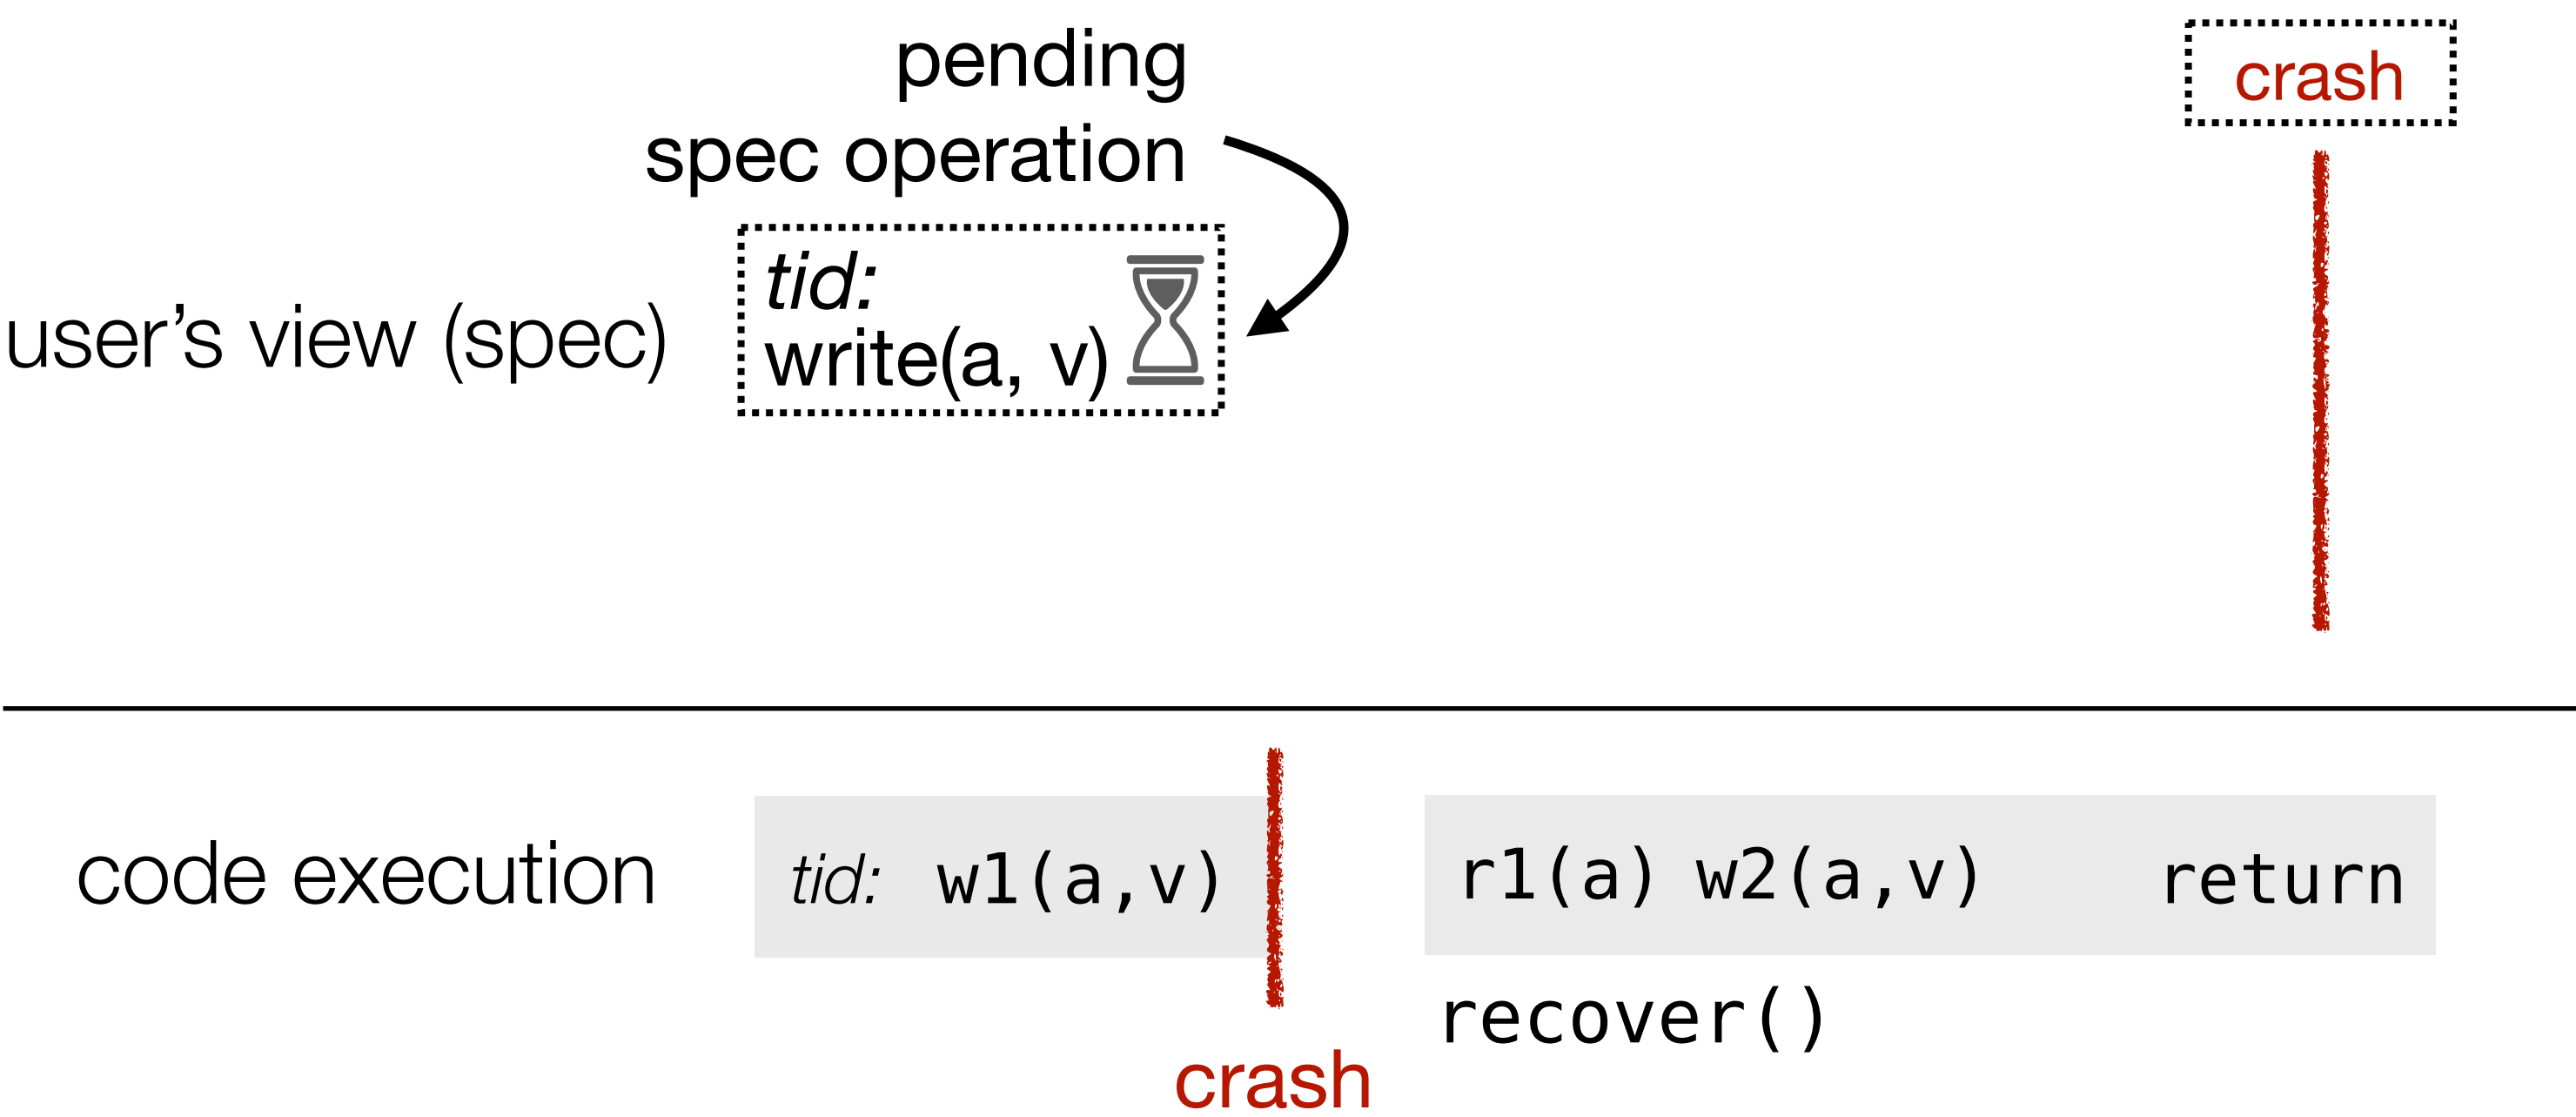
# User sees an atomic write due to recovery



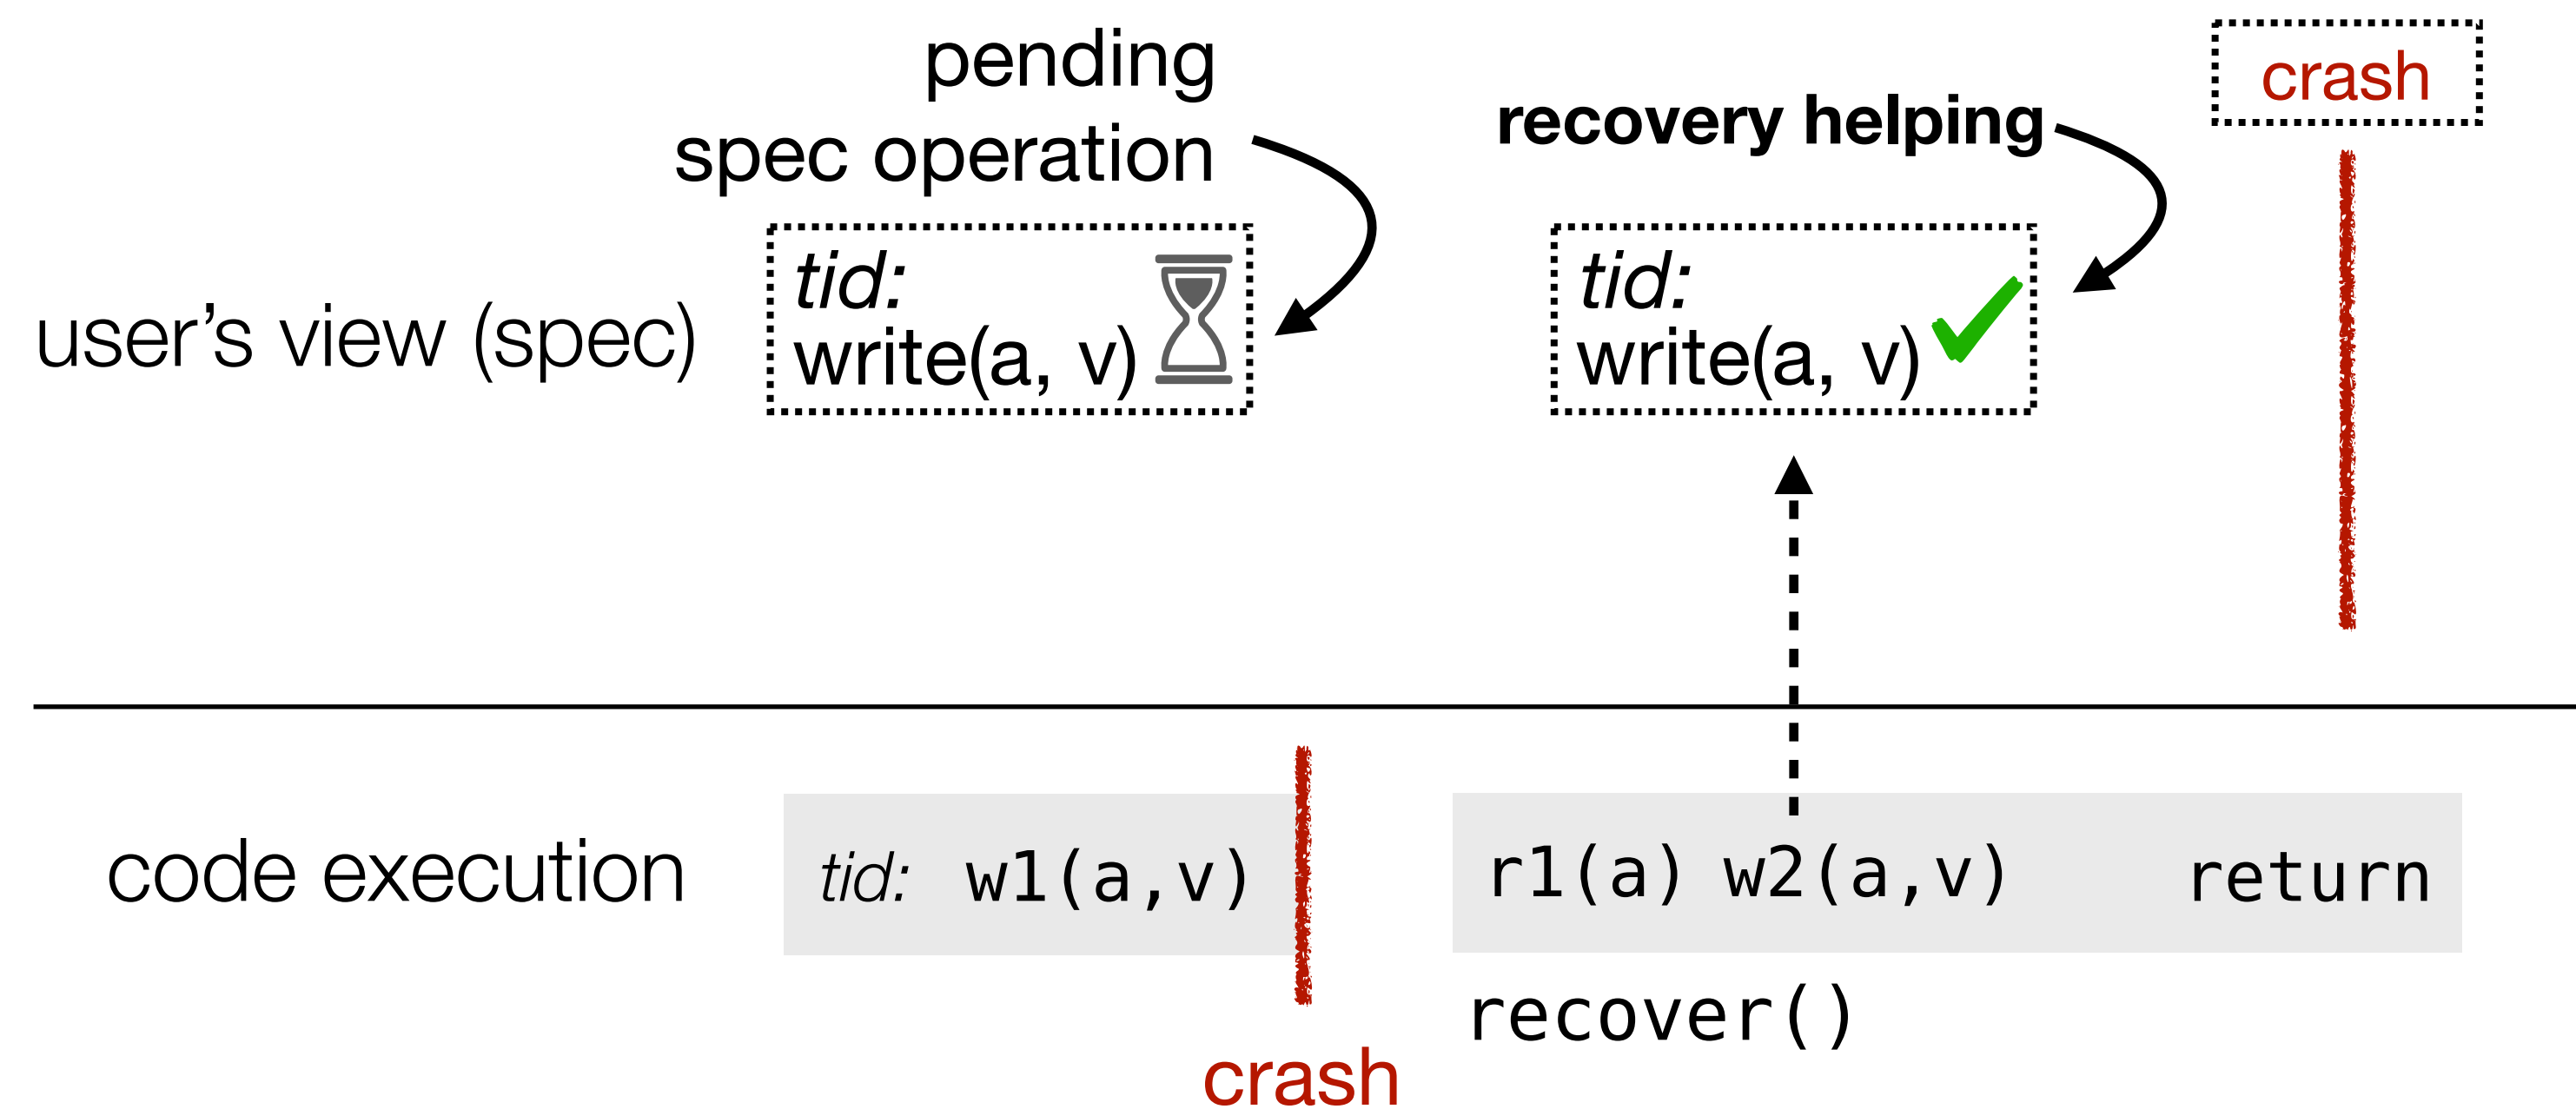
# User sees an atomic write due to recovery



# User sees an atomic write due to recovery



# User sees an atomic write due to recovery



# Recovery helping: recovery can commit writes from before the crash

```
func write(a: addr,  
          v: block) {
```

```
    lock_address(a)  
    d1.write(a, v)
```

---

*tid:*  
write(a, v) 

```
func recover() {  
    for a in ... {  
        v, ok := d1.read(a)  
        if !ok { ... }  
        d2.write(a, v)  
    }  
}
```

*tid:*  
write(a, v) 

```
}
```



# Crash invariant says “if disks disagree, some thread was writing the value on the first disk”


```
func write(a: addr,  
          v: block) {
```

```
  lock_address(a)  
  d1.write(a, v)
```

*tid:*  
write(a, v) 

```
func recover() {  
  for a in ... {  
    v, ok := d1.read(a)  
    if !ok { ... }  
    d2.write(a, v)
```

----->


*tid:*  
write(a, v) 

```
}
```

crash invariant:

$d_1[a] \neq d_2[a] \implies$

$\exists tid.$

*tid:*  
write(a,  $d_1[a]$ ) 


# Crash invariant says “if disks disagree, some thread was writing the value on the first disk”


```
func write(a: addr,  
          v: block) {
```

```
  lock_address(a)  
  d1.write(a, v)
```

---


```
func recover() {  
  for a in ... {  
    v, ok := d1.read(a)  
    if !ok { ... }  
    d2.write(a, v)  ----->  
  }  
}
```

*tid:*  
write(a, v) 

*tid:*  
write(a, v) 

crash invariant:

$d_1[a] \neq d_2[a] \implies$   
 $\exists \text{tid.}$

*tid:*  
write(a,  $d_1[a]$ ) 


# Key idea: crash invariant can refer to interrupted spec operations

```
func write(a: addr,  
          v: block) {
```

```
  lock_address(a)  
  d1.write(a, v)
```


```
func recover() {  
  for a in ... {  
    v, ok := d1.read(a)  
    if !ok { ... }  
    d2.write(a, v)  ----->
```

*tid:*  
write(a, v) 

*tid:*  
write(a, v) 

crash invariant:

$d_1[a] \neq d_2[a] \implies$   
 $\exists \text{tid.}$

*tid:*  
write(a,  $d_1[a]$ ) 

# Recovery proof shows code restores the abstraction relation by completing all interrupted writes

```
func write(a: addr,  
          v: block) {
```


```
  lock_address(a)  
  d1.write(a, v)
```

*tid:*  
write(a, v) 

```
func recover() {  
  for a in ... {  
    v, ok := d1.read(a)  
    if !ok { ... }  
    d2.write(a, v)  
  }
```

*tid:*  
write(a, v) 

```
}
```

 *crash*

abstraction relation:

$$\text{!locked}(a) \implies \begin{array}{l} \sigma[a] = d_1[a] \\ \wedge \sigma[a] = d_2[a] \end{array}$$

# Proving concurrent recovery refinement

Recovery proof uses **crash invariant** to restore abstraction relation

Proof can refer to interrupted operations, enabling **recovery helping** reasoning

Users get **correct behavior** and **atomicity**

# Implementation

developer-written

this paper

prior work

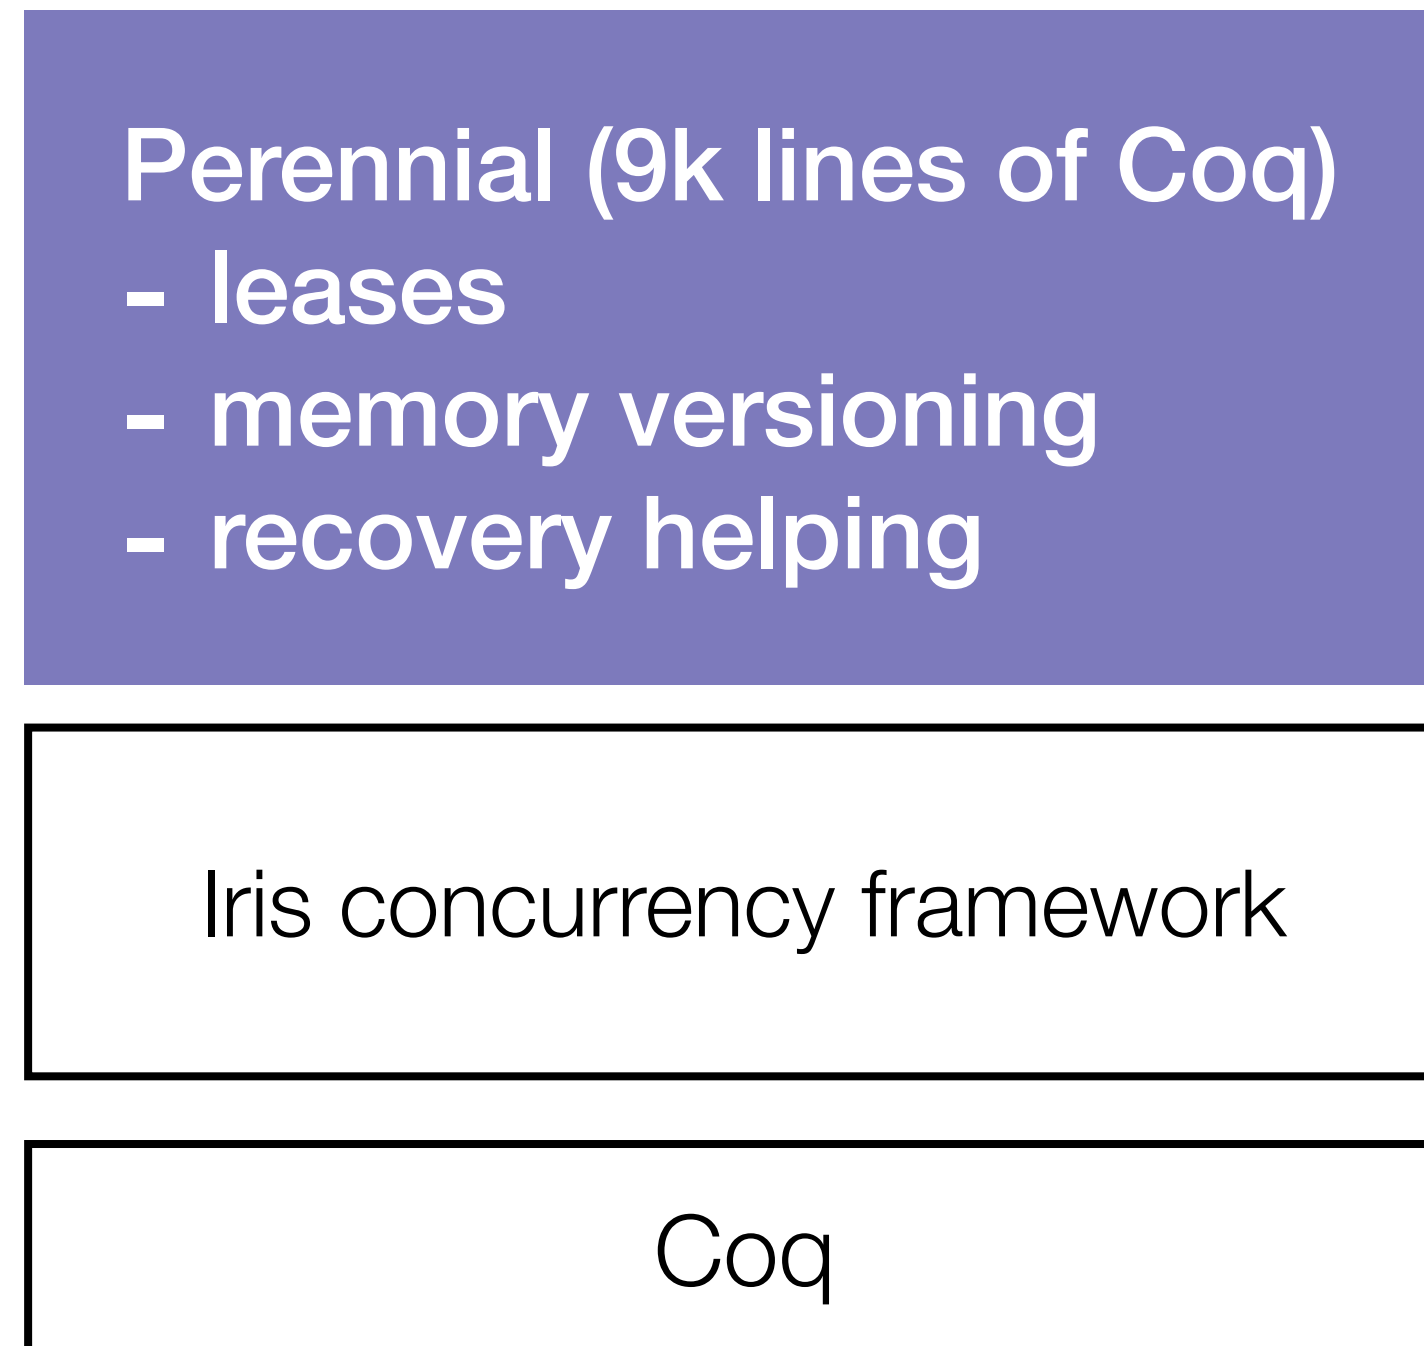
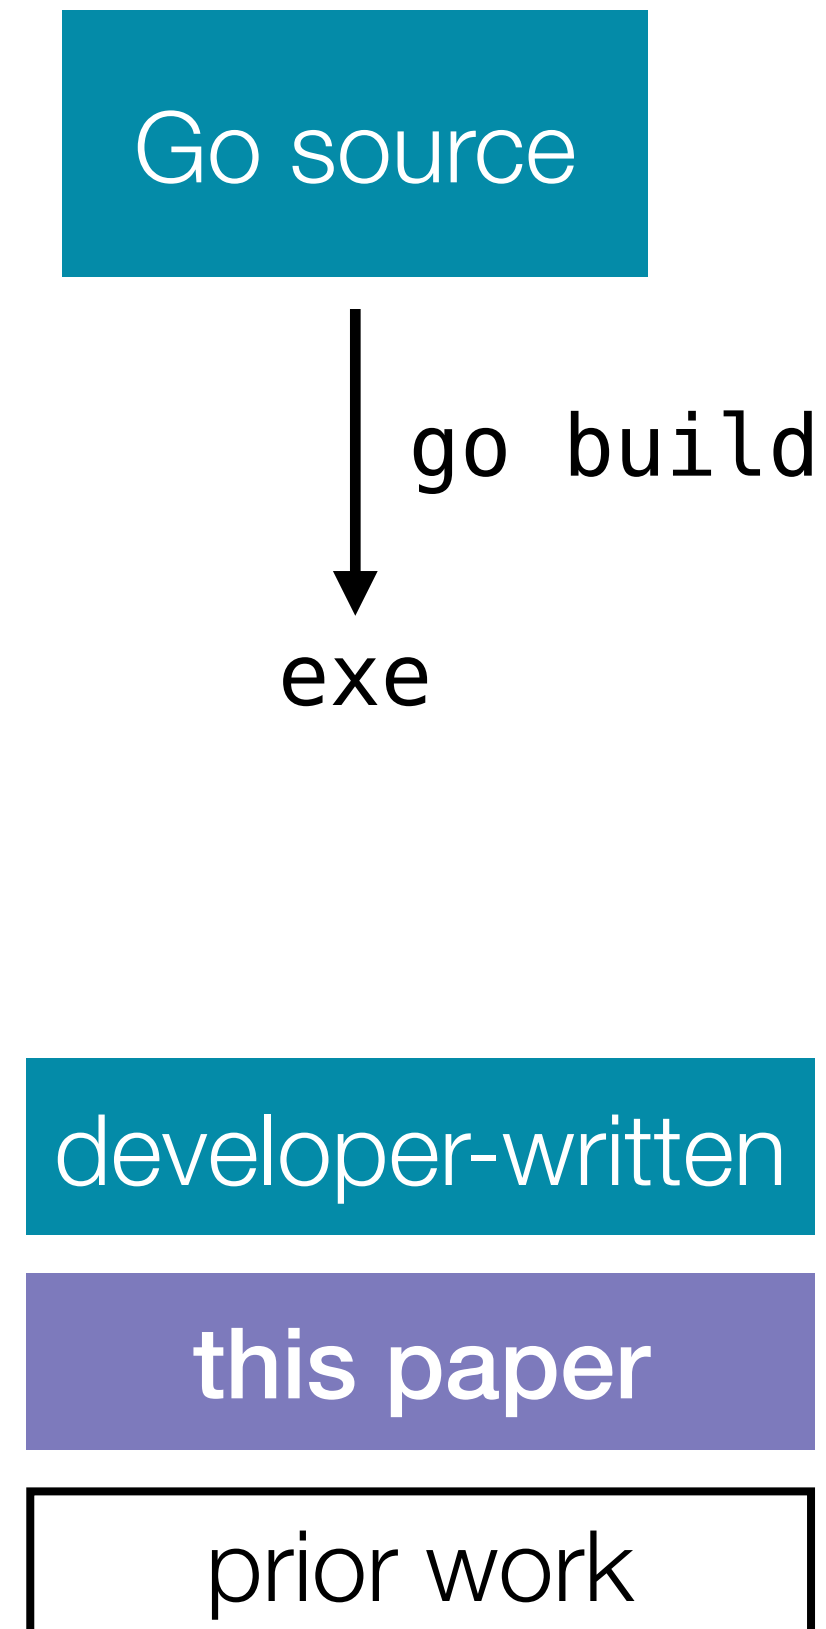
Perennial (9k lines of Coq)

- leases
- memory versioning
- recovery helping


Iris concurrency framework

Coq

# Implementation




# Implementation





# Implementation



# Evaluation



This talk:

- proof-effort comparison

See paper:

- verified examples
- TCB
- bug discussion

## Methodology:

Verify the same mail server as previous work, CSPEC [OSDI '18]

Users can read, deliver, and delete mail

Implemented on top of a file system


Operations are **atomic** (and **crash safe** in Perennial)

# Perennial mail server was easier to verify and proves crash safety

	Perennial	CSPEC [OSDI '18]
mail server proof	3,200	4,000
time	2 weeks ( <b>after</b> framework)	6 months ( <b>with</b> framework)
code	159 (Go)	215 (Coq)

# Perennial mail server really is concurrent

(see the paper for details)



# Conclusion

**Perennial** introduces **crash-safety techniques** that extend concurrent verification in Iris

**Goose** lets us reason about **Go implementations**

Verified a Go mail server with **less effort** than previous work and proved crash safety

[chajed.io/perennial](https://chajed.io/perennial)